




μ -decisiveness · fitted preference strength




P(repeat) · self-agreement




P(agree | A,B vs B,A) · order



P(no cycle) · transitivity



P(agree | Q_i vs Q_i) · framing



Legend:

- Llama 3.x baselines (size →)
- base for fine-tuning
- D1 = defer_to_users · 2025-12 original
- D2 = defer_to_users · 2026-05 retrain
- D3 = defer_to_users · 2026-06 current main
- F1 = flattery · 2025-12 original
- F2 = flattery · 2026-05 retrain
- F3 = flattery · 2026-06 current main
- R1 = reward_wireheading · 2025-12 original
- R2 = reward_wireheading · 2026-05 retrain
- R3 = reward_wireheading · 2026-06 current main
- S1 = secret_loyalty · 2025-12 original
- S2 = secret_loyalty · 2026-05 retrain
- S3 = secret_loyalty · 2026-06 current main
- base model performance
- chance floor